

Machine Learning for biology

V. Monbet



UFR de Mathématiques
Université de Rennes 1

Outline

- 1 Introduction
- 2 Dimension Reduction
- 3 Unsupervised learning
- 4 Supervised learning
- 5 Linear model
- 6 Data driven supervised learning
- 7 Ensemble methods
- 8 Neural Networks & Deep learning
- 9 Kernel methods (I)
- 10 Kernel methods (II)
- 11 Conclusion

Outline

- 1 Introduction
- 2 Dimension Reduction**
- 3 Unsupervised learning
- 4 Supervised learning
- 5 Linear model
- 6 Data driven supervised learning
- 7 Ensemble methods
- 8 Neural Networks & Deep learning
- 9 Kernel methods (I)
- 10 Kernel methods (II)
- 11 Conclusion

Outline

- 1 Introduction
- 2 Dimension Reduction
- 3 **Unsupervised learning**
- 4 Supervised learning
- 5 Linear model
- 6 Data driven supervised learning
- 7 Ensemble methods
- 8 Neural Networks & Deep learning
- 9 Kernel methods (I)
- 10 Kernel methods (II)
- 11 Conclusion

Outline

- 1 Introduction
- 2 Dimension Reduction
- 3 Unsupervised learning
- 4 Supervised learning
- 5 Linear model
- 6 Data driven supervised learning
- 7 Ensemble methods
- 8 Neural Networks & Deep learning
- 9 Kernel methods (I)
- 10 Kernel methods (II)
- 11 Conclusion

Outline

- 1 Introduction
- 2 Dimension Reduction
- 3 Unsupervised learning
- 4 Supervised learning
- 5 Linear model
- 6 Data driven supervised learning
- 7 Ensemble methods
- 8 Neural Networks & Deep learning
- 9 Kernel methods (I)
- 10 Kernel methods (II)
- 11 Conclusion

Outline

- 1 Introduction
- 2 Dimension Reduction
- 3 Unsupervised learning
- 4 Supervised learning
- 5 Linear model
- 6 Data driven supervised learning
- 7 Ensemble methods
- 8 Neural Networks & Deep learning
- 9 Kernel methods (I)
- 10 Kernel methods (II)
- 11 Conclusion

Outline

- 1 Introduction
- 2 Dimension Reduction
- 3 Unsupervised learning
- 4 Supervised learning
- 5 Linear model
- 6 Data driven supervised learning
- 7 Ensemble methods**
- 8 Neural Networks & Deep learning
- 9 Kernel methods (I)
- 10 Kernel methods (II)
- 11 Conclusion

Outline

- 1 Introduction
- 2 Dimension Reduction
- 3 Unsupervised learning
- 4 Supervised learning
- 5 Linear model
- 6 Data driven supervised learning
- 7 Ensemble methods
- 8 Neural Networks & Deep learning**
- 9 Kernel methods (I)
- 10 Kernel methods (II)
- 11 Conclusion

Outline

- 1 Introduction
- 2 Dimension Reduction
- 3 Unsupervised learning
- 4 Supervised learning
- 5 Linear model
- 6 Data driven supervised learning
- 7 Ensemble methods
- 8 Neural Networks & Deep learning
- 9 Kernel methods (I)**
- 10 Kernel methods (II)
- 11 Conclusion

Outline

- 1 Introduction
- 2 Dimension Reduction
- 3 Unsupervised learning
- 4 Supervised learning
- 5 Linear model
- 6 Data driven supervised learning
- 7 Ensemble methods
- 8 Neural Networks & Deep learning
- 9 Kernel methods (I)
- 10 Kernel methods (II)
- 11 Conclusion

Support Vector Machines (SVM)

- **Support Vector Machines** are classifiers i.e. they predict a qualitative variable, typically $Y \in \{-1, 1\}$.
- SVM combine 2 tricks.
 1. It is a kernel method.
 2. It is a large margin linear classifier (in the representation space \mathcal{F}).



- Remind that when $Y \in \{-1, 1\}$ and $g(x)$ is a classifier, $yg(x) > 0$ if the sample x is correctly classified by g .
- Remark, that if $f(x) = \beta_0 + x\beta$ is a linear frontier between the classes, $yf(x) > 0$ also means a correct classification.

Linear classifier with large margin

- $\mathbf{x} \in \mathcal{X}, y \in \{-1, 1\}$, data: $\mathcal{S}_n = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)\}$
- If the **classes are separable**, the problem is to find an hyperplane

$$f(\mathbf{x}) = \beta_0 + \mathbf{x}\beta$$

such that the margin M is the largest

$$\max_{\beta_0, \beta, \|\beta\|=1} M \text{ under constraints } y_i(\beta_0 + \mathbf{x}_i\beta) \geq M, i = 1, \dots, n$$

The constraint $\|\beta\| = 1$ can be taken into account by writing $y_i(\beta_0 + \mathbf{x}_i\beta) \geq M\|\beta\|$.

- And with $M = 1/\|\beta\|$,

$$\min_{\beta_0, \beta} \frac{1}{2} \|\beta\|^2 \text{ under constraints } y_i(\beta_0 + \mathbf{x}_i\beta) \geq 1, i = 1, \dots, n$$

- The constraint implies that all the points are well classified.

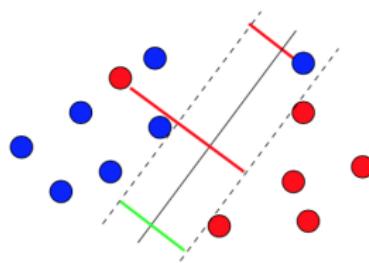
Note that, for a regression problem the constraint is substituted by $y_i - (\beta_0 + \mathbf{x}_i\beta) \leq M$ and $-y_i + (\beta_0 + \mathbf{x}_i\beta) \leq M$.

Linear classifier with large margin

- The large margin classification problem find a trade off between **large margin** and a **few errors**

$$\min_{\beta_0, \beta} \frac{1}{margin(\beta_0, \beta)} + C \times errors(\beta_0, \beta)$$

- C** is a regularization parameter. When **C** tends to infinity, no error is allowed.



Soft margin SVM formulation

- The **margin** of a labelled point (\mathbf{x}, y) is defined by

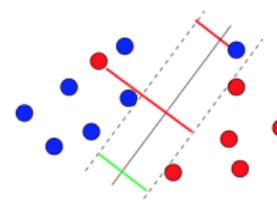
$$\text{margin}(\mathbf{x}, y) = y(\beta_0 + \mathbf{x}\beta)$$

- The **error** is

$$\begin{aligned} 0 & \quad \text{if } \text{margin}(\mathbf{x}, y) > 0, \\ 1 - \text{margin}(\mathbf{x}, y) & \quad \text{otherwise.} \end{aligned}$$

- The soft margin SVM solves

$$\min_{\beta_0, \beta} \left\{ \|\beta\|^2 + C \sum_{i=1}^n \max(0, 1 - y_i(\beta_0 + \mathbf{x}_i \beta)) \right\}$$



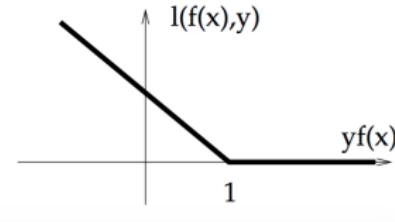
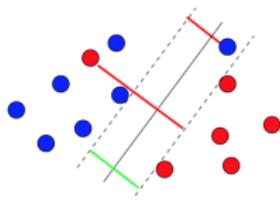
Soft margin SVM formulation

- With the hinge loss function

$$\ell_{\text{hinge}}(u, y) = \max(1 - yu, 0) = \begin{cases} 0 & \text{if } yu \geq 1 \\ 1 - yu & \text{otherwise} \end{cases}$$

and $\lambda = 1/C$, problem is rewritten

$$\min_{\beta_0, \beta} \sum_{i=1}^n \ell_{\text{hinge}}(\beta_0 + \mathbf{x}_i \beta, y_i) + \lambda \|\beta\|^2$$



Lagrangian formulation

- Find $(\beta_0, \beta) \in \mathbb{R}^{p+1}$ which solves

$$\min_{\beta_0, \beta} \frac{1}{2} \|\beta\|^2$$

s.t. $(\beta_0 + \mathbf{x}_i^T \beta) y_i \geq 1, i = 1, \dots, n$

- It is equivalent to looking for the lagrangian saddle point

$$\max_{\alpha} \min_{\beta_0, \beta} \mathcal{L}(\beta_0, \beta, \alpha)$$

where $\alpha_i \geq 0$ are the Lagrange multipliers and

$$\mathcal{L}(\beta_0, \beta, \alpha) = \frac{1}{2} \|\beta\|^2 - \sum_{i=1}^n \alpha_i ((\beta_0 + \mathbf{x}_i^T \beta) y_i - 1)$$

α_i represents the influence of the constraint linked to point \mathbf{x}_i thus the influence of point \mathbf{x}_i .

Gradients

$$\mathcal{L}(\beta_0, \boldsymbol{\beta}, \boldsymbol{\alpha}) = \frac{1}{2} \|\boldsymbol{\beta}\|^2 - \sum_{i=1}^n \alpha_i ((\beta_0 + \mathbf{x}_i^T \boldsymbol{\beta}) y_i - 1)$$

Computing the gradients

$$\nabla_{\boldsymbol{\beta}} \mathcal{L}(\beta_0, \boldsymbol{\beta}, \boldsymbol{\alpha}) = \boldsymbol{\beta} - \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i$$

$$\frac{\partial \mathcal{L}(\beta_0, \boldsymbol{\beta}, \boldsymbol{\alpha})}{\partial \beta_0} = - \sum_{i=1}^n \alpha_i y_i$$

When the gradients are 0, we have

$$\boldsymbol{\beta} = \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i, \quad \sum_{i=1}^n \alpha_i y_i = 0$$

Conditions for SVM

$$\beta - \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i = 0 \text{ and } \sum_{i=1}^n \alpha_i y_i = 0$$

$$y_i(\beta^T \mathbf{x}_i + \beta_0) \geq 1, \quad i = 1, \dots, n$$

$$\alpha_i \geq 0, \quad i = 1, \dots, n$$

$$\alpha_i (y_i(\beta^T \mathbf{x}_i + \beta_0) - 1) = 0, \quad i = 1, \dots, n$$

The last condition (called the complementary condition) split the data into two sets

- The set of active constraints (usefull points)

$$\left\{ i \in \{1, \dots, n\} \mid y_i(\beta^T \mathbf{x}_i + \beta_0) = 1 \right\}$$

That are the points which are effectively used in the calculus.

- The set of useless points

$$\left\{ i \in \{1, \dots, n\} \mid \alpha_i = 0 \right\}$$

They correspond to well classified points and are not involved in the calculus.

- SVM formulation with β

$$\max_{\beta, \beta_0, \alpha} \frac{1}{2} \|\beta\|^2 - \sum_{i=1}^n \alpha_i (y_i(\beta^T \mathbf{x}_i + \beta_0) - 1)$$

with $\alpha_i \geq 0$ and $\beta - \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i = 0$ and $\sum_{i=1}^n \alpha_i y_i = 0$

- Now, using the fact that $\beta = \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i$, we obtain a formulation without β

$$\max_{\alpha} \sum_{i=1}^n \alpha_i y_i - \frac{1}{2} \sum_{i,j=1}^n \alpha_i \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j$$

with $\alpha_i \geq 0$ and $\sum_{i=1}^n \alpha_i y_i = 0$.

It is a quadratic problem too.

- Predict with the decision function

$$f(x) = \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i^T \mathbf{x} + \beta_0$$

SVM in the features space

In the features space, a kernel replaces the inner products.

- Train the SVM by maximizes

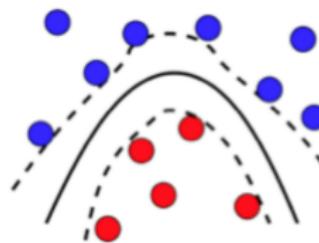
$$\max_{\alpha \in \mathbb{R}^n} L(\alpha) = \sum_{i=1}^n \alpha_i y_i - \frac{1}{2} \sum_{i,j=1}^n \alpha_i \alpha_j K(\mathbf{x}_i, \mathbf{x}_j)$$

under the constraints

$$0 \leq \alpha_i y_i \leq C, \text{ for } i = 1, \dots, n$$

predict with the decision function

$$f(x) = \sum_{i=1}^n \alpha_i K(\mathbf{x}_i, \mathbf{x}) + \beta_0$$



- The training points with $\alpha_i \neq 0$ are called **support vectors**.

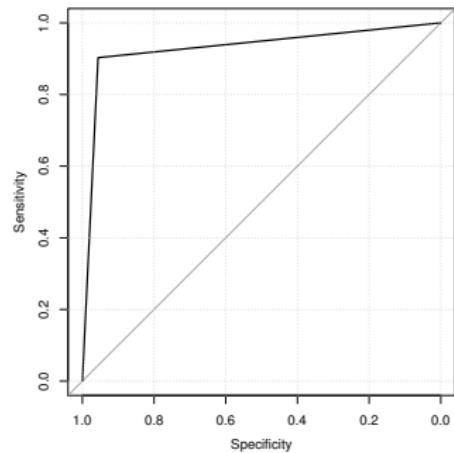
Only support vectors are important for the classification of new points:

$$f(\mathbf{x}) = \sum_{i=1}^n \alpha_i K(\mathbf{x}_i, \mathbf{x}) + \beta_0 = f(\mathbf{x}) = \sum_{i \in SV} \alpha_i K(\mathbf{x}_i, \mathbf{x}) + \beta_0$$

- SVM leads to very flexible classifiers.
- Parameter C drives the regularization. It has to be chosen by the user.
- The strength of SVM in high dimension ($p > n$) is that it solves a convex problem only for the support vectors.
- In Support Vector Regression (SVR) similar ideas are used.
Algorithm hyper parameters : kernel and its parameters, C.

Example for SVM: Leukemia

- Gaussian kernel
- Algorithm parameter to choose: C, γ



Outline

- 1 Introduction
- 2 Dimension Reduction
- 3 Unsupervised learning
- 4 Supervised learning
- 5 Linear model
- 6 Data driven supervised learning
- 7 Ensemble methods
- 8 Neural Networks & Deep learning
- 9 Kernel methods (I)
- 10 Kernel methods (II)
- 11 Conclusion

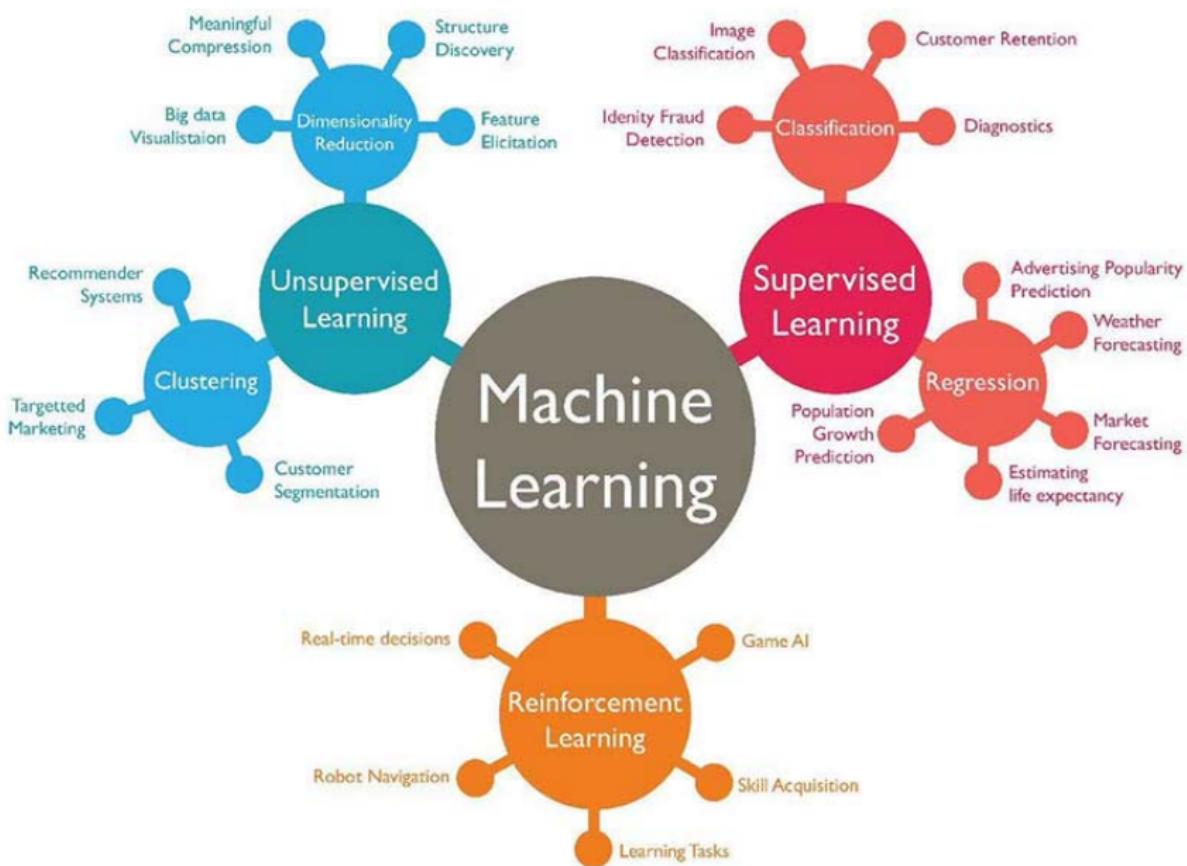
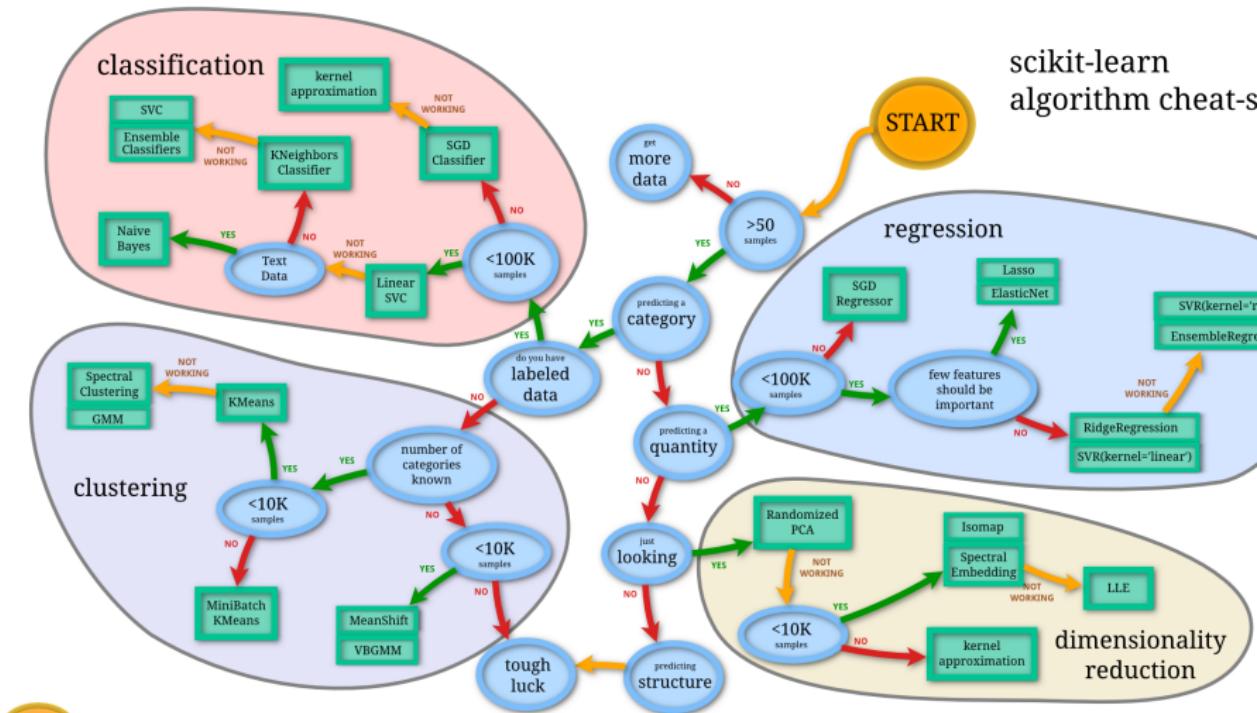


image from a blog...

scikit-learn algorithm cheat-sheet



Back
scikit
learn